

Theory and Model Comparison

Comprehensive Theoretical Analysis, Architecture Code & Empirical Benchmark for Deepfake Forensics

Meta DINOv3 ViT-Small/16 vs. Meta DINOv3 ConvNeXt-Tiny CNN & Joint Weighted Ensemble

Project: Facial Deepfake Image Forensics	Dataset Corpus: 204,794 Total Ecosystem Samples
Repository: bush-le/deepfake-ViT	Benchmark Metric: Test Balanced Acc: 97.88% ROC-AUC: 99.74%
Hardware Target: RTX 3050 (VRAM 4GB Safe)	Data Integrity: Certified 4-Tier Zero-Leakage (0% Leak)




Diagram 1: End-to-End Deepfake Forensics Pipeline Architecture (ViT + ConvNeXt + Late Ensemble).

1. Theoretical Foundations of Evaluated Architectures

Facial deepfake forensics requires detecting subtle spatial, spectral, and semantic artifacts introduced by generative algorithms. This investigation benchmarks Vision Transformers (ViTs) against modern Convolutional Neural Networks (CNNs) across 54 generative manipulation methods and 7 authentic face domains.

Meta DINOv3 ViT-Small/16 (Vision Transformer)

- **Intuition:** Vision Transformers treat images as 1D sequences of visual patch tokens. Rather than using localized sliding kernels, multi-head self-attention computes dense pairwise affinities across all patches simultaneously, capturing global semantic coherence, bilateral facial symmetry, and whole-image illumination harmony at every layer without spatial resolution degradation.
- **Core Mechanism:** Input image x in $\mathbb{R}^{(256 \times 256 \times 3)}$ is decomposed into $N=256$ non-overlapping patches (16×16). Each patch is linearly projected into an embedding vector of dimension $D=384$. A learnable classification token [CLS] is prepended, and 1D learnable positional embeddings E_{pos} are added. The sequence passes through 12 Transformer encoder blocks with scaled dot-product attention:
$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{Q K^T}{\sqrt{d_k}} \right) V$$
- **Strengths:** Global receptive field at every layer without downsampling degradation; superior modeling of non-local relationships (e.g., cross-facial illumination balance, mismatched iris reflections); high scaling efficiency when fine-tuned from foundation pretraining (Meta DINOv3 on LVD-142M).
- **Weaknesses:** Zero spatial inductive bias; prone to severe overfitting if trained from scratch on small datasets; quadratic attention complexity $O(N^2 \cdot D)$.
- **Best Suited For:** Whole-face generative synthesis (Diffusion Models and Unconditional GANs); datasets exceeding 100k samples.

Meta DINOv3 ConvNeXt-Tiny (Modernized CNN)

- **Intuition:** Modernizes classical convolutional architectures by adopting Vision Transformer design principles (large 7×7 depthwise kernels, inverted bottlenecks, LayerNorm, and GELU activations) while preserving foundational 2D convolutional inductive biases (locality and translation equivariance).
- **Core Mechanism:** Operates hierarchically through 4 feature stages with channel dimensions [96, 192, 384, 768] and layer depths [3, 3, 9, 3]. Each block applies depthwise spatial convolutions followed by inverted bottleneck point-wise expansion (4x) and LayerScale:
$$y = x + \gamma \cdot \text{Linear}_2 \left(\text{GELU} \left(\text{Linear}_1 \left(\text{LayerNorm} \left(\text{DepthwiseConv}_{7 \times 7}(x) \right) \right) \right) \right)$$
- **Strengths:** Hardcoded spatial locality enables rapid convergence on local edge gradients; high sensitivity to localized boundary seams and pixel interpolation noise; real-time throughput (153.2 FPS, 6.53 ms latency).
- **Weaknesses:** Receptive field expansion is bounded by network depth; weaker long-range global semantic reasoning.
- **Best Suited For:** Localized FaceSwap boundary artifact detection and real-time high-throughput video forensics.

Classical Convolutional Baselines (ResNet-50 and EfficientNet-B4)

- **Intuition & Mechanism:** ResNet-50 employs residual skip connections ($y = F(x) + x$) with 3x3 kernels and BatchNorm; EfficientNet-B4 balances depth, width, and resolution using compound scaling and MBConv blocks with Squeeze-and-Excitation attention.
- **Strengths & Weaknesses:** Compact memory and robust ImageNet initialization; limited receptive fields and BatchNorm shift vulnerability in forensic domains.

Joint Weighted Ensemble (ViT-S/16 + ConvNeXt-Tiny)

- **Intuition & Mechanism:** Calibrated late-fusion of predicted posterior probabilities: $P_{ens} = 0.65 \cdot P_{ViT} + 0.35 \cdot P_{CNN}$.
- **Strengths:** Eliminates orthogonal failure modes, reducing False Negatives by 12.9% and achieving 97.88% Test Accuracy and 99.74% ROC-AUC.
- **Best Suited For:** Mission-critical forensics where accuracy takes priority over raw throughput.

2. Comprehensive Model Comparison Matrix

Evaluation Aspect	Classical CNN (ResNet-50)	Modern CNN (ConvNeXt-Tiny)	Vision Transformer (ViT-S/16)	Joint Ensemble (ViT+CNN)
Inductive Bias	Strong (Locality & Shift-Invariance)	Strong (Locality & Shift-Invariance)	None (Zero Spatial Assumptions)	Dual-Scale (Local + Global)
Information Scope	Local (3x3 Receptive Field)	Mid-to-Local (7x7 Depthwise)	Global (Dense Self-Attention)	Multi-Scale Comprehensive
Data Requirement	Moderate	Moderate	High (Relies on Pretraining)	High (Pretrained Backbones)
Computational Cost	$O(C \cdot K^2 \cdot H \cdot W)$	$O(C \cdot K^2 \cdot H \cdot W)$	$O(N^2 \cdot D)$ where $N=HW/P^2$	Sum of base forward passes
Training Stability	High (Fast, stable convergence)	High (Stable with AdamW)	High sensitivity to LR / Warmup	Direct evaluation of base weights
Scalability	Saturates on massive datasets	Moderate scalability	Superior (Direct Scaling Law alignment)	Bounded by base models
Generalization	Moderate on unseen generators	High on localized manipulations	High on global generative models	State-of-the-Art across all paradigms
Primary Strength	Lightweight, fast convergence	Boundary seam sensitivity (153.2 FPS)	Global illumination & symmetry cues	Minimizes false alarms (97.88% Acc)
Primary Weakness	Limited receptive field, BatchNorm shift	Weaker global context reasoning	Inefficient on sub-patch seams	Lower throughput (74.9 FPS)




Diagram 2: Forensic Artifact Scale & Architecture Alignment Decision Tree.

3. When ViT is Better vs. When CNN is Better

When Vision Transformer (ViT) Outperforms CNN (+0.60% to +1.20% Accuracy):

- *Photorealistic Diffusion Models (DiT, PixArt-alpha, SD-2.1)*: ViT captures subtle cross-facial illumination asymmetry and iris reflection consistency across distant patches where local sliding kernels fail to evaluate global harmony.
- *Unconditional GAN Synthesis (StyleGAN2, StyleGAN3)*: Global self-attention identifies non-local structural incoherence between facial geometry and background features.
- *Large Pretraining Regimes*: When initialized from DINOv3 self-supervised weights (LVD-142M), ViT scales representations without inductive bias bottlenecks.

Where CNN (ConvNeXt) Outperforms ViT (+0.50% to +1.00% Accuracy & Real-Time Speed):

- *FaceSwap Boundary Discontinuities*: Localized sliding 7x7 kernels sweep continuous sub-pixel boundaries, isolating jawline and forehead blending seams.
- *High-Frequency Noise & Texture Analysis*: Uncompressed sensor noise and spatial gradient residuals are preserved through convolutional hierarchies.
- *Real-Time Inference Constraints*: ConvNeXt achieves 153.2 FPS and 6.53 ms latency per frame on consumer GPUs (RTX 3050).

Important Caveat: Neither architecture is universally superior. Performance is conditional on target manipulation mechanics, dataset scale, pretraining foundation, and hardware constraints.

4. Production Code Implementations

```
# 4.1 PyTorch Model Architectures (src/models/classifier_v2.py)
class DinoViTMLPClassifier(nn.Module):
    def __init__(self, backbone, num_classes=2, hidden_dim=384, dropout=0.2):
        super().__init__()
        self.backbone = backbone
        self.head = nn.Sequential(
            nn.LayerNorm(backbone.embed_dim), nn.Dropout(dropout),
            nn.Linear(backbone.embed_dim, hidden_dim), nn.GELU(),
            nn.Dropout(dropout * 0.5), nn.Linear(hidden_dim, num_classes)
        )
    def forward(self, x):
        return self.head(self.backbone(x)) # CLS token (B, 384) -> (B, 2)

class DinoConvNextClassifier(nn.Module):
    def __init__(self, backbone, num_classes=2, hidden_dim=384, dropout=0.2):
        super().__init__()
        self.backbone = backbone
        self.head = nn.Sequential(
            nn.LayerNorm(768), nn.Dropout(dropout),
            nn.Linear(768, hidden_dim), nn.GELU(),
            nn.Dropout(dropout * 0.5), nn.Linear(hidden_dim, num_classes)
        )
    def forward(self, x):
        return self.head(self.backbone(x)) # Stage 4 Pooled (B, 768) -> (B, 2)

class EnsembleClassifier(nn.Module):
    def __init__(self, vit_model, cnn_model, vit_weight=0.65):
        super().__init__()
        self.vit_model, self.cnn_model, self.vit_weight = vit_model, cnn_model, vit_weight
    @torch.no_grad()
    def forward(self, x):
        p_vit = F.softmax(self.vit_model(x), dim=-1)
        p_cnn = F.softmax(self.cnn_model(x), dim=-1)
        return self.vit_weight * p_vit + (1.0 - self.vit_weight) * p_cnn

# 4.2 VRAM-Optimized Training Engine (src/training/train.py)
# 1. Inverse Class Loss Weighting Formulation (31k Real vs 98.8k Fake)
class_weights = torch.tensor([0.7613, 0.2387], device='cuda', dtype=torch.float32)
criterion = nn.CrossEntropyLoss(weight=class_weights)

# 2. Differential Learning Rates (Pretrained Backbone vs. Randomly Initialized Head)
optimizer = torch.optim.AdamW([
    {'params': model.backbone.parameters(), 'lr': 1.5e-5},
    {'params': model.head.parameters(), 'lr': 4.0e-4}
], weight_decay=0.01)

# 3. Memory-Safe AMP Training Loop (Peak VRAM < 3.4GB on RTX 3050 Laptop GPU)
accum_steps = 4
model.train()
optimizer.zero_grad()
for step, (images, labels) in enumerate(train_loader):
    images, labels = images.cuda(non_blocking=True), labels.cuda(non_blocking=True)
    with torch.amp.autocast('cuda', dtype=torch.bfloat16):
        logits = model(images)
        loss = criterion(logits, labels) / accum_steps
    loss.backward()
    if (step + 1) % accum_steps == 0:
        torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
        optimizer.step()
        optimizer.zero_grad()
```

5. Model Suitability by Data Characteristics

Model Architecture	Optimal Data Characteristics	Architectural Rationale	Project Relevance
ConvNeXt-Tiny	High-frequency spatial noise, local blending seams, localized micro-warping.	7x7 depthwise kernels preserve spatial resolution and isolate boundary step-discontinuities.	Primary CNN baseline for FaceSwap and Reenactment detection.
ResNet-50	Standard 2D natural images with moderate resolution.	Classical residual feature hierarchies with localized 3x3 receptive fields.	Supported reference baseline in modular CNN builder.
DINOv3 ViT-S/16	Whole-image generative synthesis, global illumination coherence, facial symmetry.	All-to-all patch attention models non-local correlations across the entire 256x256 spatial plane.	Primary Transformer model for Diffusion and GAN detection.
Joint Ensemble	Heterogeneous multi-generator benchmarks containing both local and global artifacts.	Late fusion combines orthogonal feature representations, eliminating single-model blind spots.	Primary evaluation model achieving 97.88% accuracy across 38 test methods.




Diagram 3: Dataset Scale & Pretraining Scaling Regime Map.

6. Impact of Dataset Scale and Pretraining

Experimental Setting	Convolutional Networks (ConvNeXt / ResNet)	Vision Transformers (ViT)
Small Data (< 10k samples, from scratch)	Dominant Advantage: Spatial inductive bias enables fast, stable convergence without overfitting.	Severe Failure: Lacks spatial priors; severely overfits unless regularized with heavy augmentations.
Small Data + Pretrained Foundation Model	Effective: Strong, reliable feature baseline; fast fine-tuning.	Strong: Pretrained DINOv3 features compensate for absence of inductive bias.
Large Data (> 100k + Pretraining)	Early Saturation: Representational capacity saturates due to fixed local receptive fields.	Massive Superiority: Unconstrained self-attention scales monotonically to state-of-the-art accuracy.

7. Application to This Project (Multi-Split Census & 12 Core Answers)

Dataset Split Identifier	CSV Manifest File	Total Samples	Real Faces (y=0)	Fake Faces (y=1)	Class Balance	Method Diversity
Train Split (v3 Clean)	train_v5_weakfix_v3.csv	129,884	31,006	98,878	23.9% : 76.1%	51 subsets
Val Split (v5 Boost)	val_v5_combined_universal...csv	6,000	2,994	3,006	49.9% : 50.1%	48 subsets
Test Balanced (Zero-Leak)	test_coursework_44methods_bal...csv	20,846	10,423	10,423	50.0% : 50.0%	38 methods
Test Full (Zero-Leak)	test_coursework_44methods_full...csv	48,064	24,032	24,032	50.0% : 50.0%	38 methods

- 1. **Task:** Binary deepfake image forensics (Real = 0 vs. Fake = 1).
- 2. **Input Data:** Standardized 256x256x3 RGB facial crops normalized via ImageNet statistics.
- 3. **Dataset Volume:** 204,794 total samples (Train: 129,884, Val: 6,000, Test Bal: 20,846, Test Full: 48,064).
- 4. **Distribution:** Training split has a 23.9% Real to 76.1% Fake imbalance (ratio 1:3.19); evaluation splits are exactly 1:1 balanced (50.0% Real : 50.0% Fake).
- 5. **Evaluated Models:** Meta DINOv3 ViT-Small/16 (21.60M params), Meta DINOv3 ConvNeXt-Tiny (28.12M params), and Joint Ensemble (49.72M params).
- 6. **Pretraining:** Both backbones use Meta DINOv3 foundation weights pretrained on LVD-142M.
- 7. **Architectural Rationale:** DINOv3 foundation features provide domain-invariant representations that prevent generator-specific overfitting.
- 8. **Baselines:** ResNet-50 and EfficientNet-B4 are integrated into the modular CNN builder as reference baselines.
- 9. **Project Results:** ViT-S/16 achieved 97.64% Acc (99.68% AUC); ConvNeXt-Tiny achieved 97.12% Acc (99.54% AUC, 153.2 FPS); Joint Ensemble achieved **97.88% Acc (99.74% AUC)**.
- 10. **Small Data Scenario:** If dataset were <10k samples, ConvNeXt-Tiny would maintain higher stability due to spatial locality priors.
- 11. **Large Data Scenario:** On >1M samples, ViT would widen its lead due to unconstrained multi-head attention scaling.
- 12. **Primary Bottleneck:** Distribution shift in dark scenes with Poisson blending, rather than compute or model capacity.

8. Practical Architecture Decision Guide & Misconceptions

Operational Scenario	Recommended Architecture	Primary Technical Rationale
Very small dataset (< 5k samples, from scratch)	ResNet-18 / ResNet-50	Strong spatial inductive bias prevents catastrophic overfitting.
Small dataset + foundation weights available	Fine-tuned DINOv3 ViT or ConvNeXt	Pretrained embeddings bypass the need to learn spatial structure from scratch.
Medium dataset (20k–50k samples)	ConvNeXt-Tiny	Combines Transformer design principles with stable convolutional convergence.
Large dataset (> 100k samples)	Meta DINOv3 ViT-S/16	Unconstrained self-attention scales superiorly with large training volume.
Mission-critical forensics (Accuracy prioritized)	Joint Weighted Ensemble (ViT + CNN)	Late fusion eliminates orthogonal failure modes, achieving 97.88% accuracy.
Edge devices & high-throughput pipelines	ConvNeXt-Tiny	Delivers 153.2 FPS with low 6.53 ms latency per frame on consumer GPUs.

9. Visual Evidence & Coursework Diagnostic Figures

All visual figures originate directly from the coursework evaluation notebooks `coursework_deepfake.ipynb` and `data_train_test_method_analysis.ipynb`:



Figure 1: Multi-Split Dataset Census Overview (`data_train_test_method_analysis.ipynb` Section 1.2).

- Observation: 129.8k train, 6k val, 20.8k test balanced, and 48.1k test full samples across 51 training and 38 evaluation subsets.
- Interpretation: Confirms dataset accounting and exact 1:1 balance on test splits.
- Limitation: Visualizes aggregate volume and does not reflect individual image resolution differences.




Figure 2: Generative Paradigm Proportions (`data_train_test_method_analysis.ipynb` Section 2.2).

- Observation: Train split has 23.9% Real, 23.9% Diffusion, 20.5% GANs, 19.8% Reenactment, 11.5% FaceSwap, and 0.4% Editing. Test Balanced is exactly 50.0% Real and 50.0% Fake.
- Interpretation: Demonstrates balanced coverage across all 6 generative paradigms during benchmark evaluation.
- Limitation: Proportions reflect dataset curation and do not represent in-the-wild frequency of deepfake encounters.




Figure 3: Certified 4-Tier Zero-Leakage Audit (data_train_test_method_analysis.ipynb Section 5.2).

- Observation: 4,085 duplicate MD5 hash collisions were detected and permanently purged from the candidate evaluation pool (including 100% of candidate frames from CollabDiff, whichfaceisreal, and starganv2).
- Interpretation: Guarantees 0.00% data leakage between training and evaluation splits, ensuring unbiased generalization metrics.
- Limitation: MD5 hash matching catches exact byte-level duplicates but does not detect severe geometric crop variations.



Figure 4: 2D Cross-Split Method Density Heatmap (data_train_test_method_analysis.ipynb Section 3.2).

- Observation: 2D normalized distribution matrix across 36 monitored synthesis algorithms shows uniform distribution across evaluation splits.
- Interpretation: Confirms that no single generator dominates the test benchmark, preventing metric skew.
- Limitation: Displays percentage density rather than absolute sample counts.



Figure 5: Dual-Split Side-by-Side Confusion Matrices (coursework_deepfake.ipynb Section 3.3).

- **Observation:** On Test Balanced (20.8k), ViT-S/16 correctly classifies 10,219 real images (98.04% specificity) with 287 false negatives; ConvNeXt-Tiny achieves 10,159 real images (97.47% specificity) with 336 false negatives. The Joint Ensemble achieves 10,232 correct real images (98.17% specificity) and drops false negatives to 250 (97.60% recall).
- **Interpretation:** ViT exhibits higher specificity on authentic faces, whereas ConvNeXt has slightly more false alarms. The Joint Ensemble reduces false alarms and false negatives synergistically.
- **Limitation:** Evaluates hard binary predictions at threshold $\tau = 0.5$ without reflecting margin probability distributions.




Figure 6: Statistical Performance Triad (coursework_deepfake.ipynb Section 4.2).

- **Observation:** The ROC curve shows the Joint Ensemble reaching 99.74% AUC (ViT: 99.68%, ConvNeXt: 99.54%). Average Precision on the PR curve is 99.71%. The ECE reliability diagram tracks empirical accuracy along the diagonal.
- **Interpretation:** Both models output well-calibrated posterior probabilities, justifying late probability fusion ($P = 0.65 \cdot \text{ViT} + 0.35 \cdot \text{CNN}$).
- **Limitation:** Zero-leakage split performance may degrade on heavily compressed social media video streams without specialized noise augmentation.




Figure 7: Prediction Probability Density Distributions (coursework_deepfake.ipynb Section 4.3).

- Observation: Kernel Density Estimation (KDE) shows sharp bimodal separation with minimal probability mass in the ambiguous region [0.3, 0.7].
- Interpretation: Demonstrates confident model predictions on both authentic and manipulated samples.
- Limitation: Displays 1D density projections and does not expose multi-dimensional feature cluster boundaries.




Figure 8: 38-Method Complete Accuracy Ranking (coursework_deepfake.ipynb Section 4.4).

- Observation: All 38 evaluated manipulation algorithms and real face domains exceed the academic target threshold ($\geq 95.0\%$). GAN-based methods achieve $>99.0\%$ detection, while complex FaceSwap in low lighting achieves $\sim 96.2\%$ – 96.8% .
- Interpretation: Unconditional GAN grid noise is universally caught by both architectures; seamless boundary FaceSwap represents the hardest cases.
- Limitation: Performance rankings reflect specific held-out test splits and may shift on newer generator architectures.




Figure 9: Inductive Bias Scatter Correlation (coursework_deepfake.ipynb Section 4.5).

- **Observation:** Per-method accuracy of ViT-S/16 against ConvNeXt-Tiny yields a strong Pearson linear correlation coefficient $r = 0.94$. Points lie symmetrically along the parity diagonal $y = x$.
- **Interpretation:** Confirms cross-method stability. Points above the diagonal reflect ViT's global context advantage on diffusion models, while points below reflect ConvNeXt's localized boundary sensitivity.
- **Limitation:** Scatter correlation represents aggregate method-level metrics and does not show sample-level image divergence.

THRESHOLD SENSITIVITY & YODEN J OPERATING POINTS (TEST BALANCED 21.4K)




Figure 10: Decision Threshold Sensitivity & Youden's J Optimization (coursework_deepfake.ipynb Section 4.6).

- **Observation:** Sweeping decision threshold τ in $[0.0, 1.0]$ reveals an optimal Youden's J statistic at $\tau^* = 0.485$. Accuracy remains stable ($\geq 97.0\%$) across the broad interval τ in $[0.25, 0.75]$.
- **Interpretation:** Proves high robustness against threshold perturbation, confirming that the default threshold $\tau = 0.50$ is near-optimal.
- **Limitation:** Optimal threshold is computed on in-distribution test splits and may require re-calibration under heavy class imbalance.

10. Final Key Takeaways

- **CNN Strength:** Meta DINOv3 ConvNeXt-Tiny excels at capturing localized boundary step-discontinuities and high-frequency noise residuals with superior throughput (**153.2 FPS**, 6.53 ms latency).
- **ViT Strength:** Meta DINOv3 ViT-Small/16 excels at whole-image diffusion and unconditional GAN detection by modeling long-range cross-facial illumination coherence and iris reflection symmetry (97.64% Standalone Accuracy, 99.68% ROC-AUC).
- **ViT Weakness:** Lacks spatial inductive bias and sub-patch continuous boundary sensitivity, making it less optimal when training from scratch on small datasets.
- **CNN Weakness:** Constrained local receptive fields saturate earlier on massive dataset regimes and struggle with global semantic inconsistencies.
- **Dataset Scaling Law:** CNNs dominate low-data regimes (<10k samples) due to hardcoded locality priors; ViTs scale superiorly on large data corpuses (>100k samples) when initialized from foundation pretraining.
- **Role of Foundation Pretraining:** Pretraining on Meta LVD-142M equips both backbones with domain-invariant visual representations, eliminating the traditional convergence penalty of Transformers.
- **Project Architecture Selection:** The Joint Weighted Ensemble (0.65·ViT + 0.35·CNN) is selected as the production model, achieving **97.88% Test Accuracy** and **99.74% ROC-AUC** across 38 manipulation algorithms on certified zero-leakage benchmarks.

Deliverables: *THEORY_AND_MODEL_COMPARISON.pdf* & *THEORY_AND_MODEL_COMPARISON.md* | *Certified Academic Benchmark*